

CartPole

January 17, 2020

```
In [ ]: import gym
import os
import time
import random
import numpy as np
from collections import deque
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam
import matplotlib.pyplot as plt
env = gym.make('CartPole-v0')
obs = env.reset()
state_size = env.observation_space.shape[0]
action_size = env.action_space.n
batch_size = 32
n_episodes = 1001
output_dir = 'model_output/cartpole'
if not os.path.exists(output_dir):
    os.makedirs(output_dir)

# Define agent

class DQNAgent:
    def __init__(self, state_size, action_size):
        self.state_size = state_size
        self.action_size = action_size
        self.memory = deque(maxlen = 2000)
        self.gamma = 0.95
        self.epsilon = 1.0
        self.epsilon_decay = 0.995
        self.epsilon_min = 0.01
        self.learning_rate = 0.1

        self.model = self._build_model()

    def _build_model(self):
        model = Sequential()
```

```

        model.add(Dense(32,input_dim = self.state_size, activation = 'relu'))
        model.add(Dense(32,activation = 'relu'))
        model.add(Dense(self.action_size, activation = 'linear'))
        model.compile(loss = 'mse', optimizer = Adam(lr = self.learning_rate))
        return model

    def remember(self, state, action, reward, next_state, done):
        self.memory.append((state, action, reward, next_state, done))

    def act(self, state):
        if np.random.rand() <= self.epsilon:
            return random.randrange(self.action_size)
        act_values = self.model.predict(state)
        return np.argmax(act_values[0])

    def replay(self, batch_size):
        minibatch = random.sample(self.memory, batch_size)

        for state, action, reward, next_state, done in minibatch:
            target = reward
            if not done:
                target = reward + self.gamma * np.max(self.model.predict(next_state)[0],
                                                       target_f = self.model.predict(state))
            target_f[0][action] = target

            self.model.fit(state, target_f, epochs = 1, verbose = 0)

            if self.epsilon > self.epsilon_min:
                self.epsilon *= self.epsilon_decay

    def load(self, name):
        self.model.load_weights(name)
    def save(self, name):
        self.model.save_weights(name)

agent = DQNAgent(state_size, action_size)

# interact with the environment

done = False

for e in range(n_episodes):

    state = env.reset()
    state = np.reshape(state, [1, state_size])

```

```

env.render()

action = agent.act(state)
next_state, reward, done, _ = env.step(action)
reward = reward if not done else -100
next_state = np.reshape(next_state, [1, state_size])
agent.remember(state, action, reward, next_state, done)
state = next_state

if done:
    print('episode: {}/{}, score:{}, e:{:.2}'.format(e, n_episodes, time, agent.epsi))
    break
env.close()

if len(agent.memory) > batch_size:
    agent.replay(batch_size)
if e % 50 == 0:
    agent.save(output_dir + 'weights_' + '{:4d}'.format(e) + '.hdf5')

```